

Let us (humans, broadly conceived) just [understand more/better]/[become more skillful/capable] (instead of making a distinct artifact more intelligent/capable than us), maybe forever

July 2024

Abstract

It is (in my present view) likely that if we humans continue to proceed around artificial intelligence as carelessly as we have been so far, we will — probably in the next century, and plausibly in the next few decades — create and [(largely unwillingly) concede control of the future to] an artificial system more intelligent/capable/powerful than humanity, which won't be anywhere close to being “our worthy successor”. I would like us to pick another path that improves on this (in my present view) profoundly stupid default path. The aim of this essay is to propound an alternative to (maybe ever) creating a distinct artificial system which surpasses our abilities: for humanity to indefinitely continue understanding more and becoming more skillful.

list of contents:

- spelling out a distinction between us (humanity) becoming smarter vs creating a distinct entity smarter than us
- how ‘let’s become smarter’ is distinct from other visions for a worthwhile future containing beings much more intelligent than we are now. us becoming smarter indefinitely looks very different from the futures most present technical alignment researchers mostly aim for — it looks very different from what most presently have in mind when talking about ‘solving the alignment problem’. These futures involve creating another mind much smarter than us which is distinct from us but ‘serves our values’; this obviously includes an ‘AGI sovereign’ but also includes ‘corrigible AGI’ (these are two canonical options, e.g. figuring prominently in Yudkowsky 2022. the future
- a section arguing that any other plan around stuff becoming much smarter working would be highly anti-natural or something. this would comparisons with stuff like an AI sovereign or corrigible AI. but what about trying to use AI to solve uploading?
- which pursuits currently seem promising for making the future such?
 - broadly, making it so humanity can better-choose its future
 - conceiving of and building systems of governance which make it possible for humanity to choose its future
 - making it so humanity has the option to prevent the construction of an artificial intelligence smarter than humanity
 - making it so we choose to execute that option (tentatively)
 - in some sense, anything that makes humans more intelligent and skillful counts. so any usual philosophy, mathematics, science matters here to some extent
 - but especially work which builds a clearer understanding of our position and makes us more able to choose a future for ourselves — so perhaps especially a particular kind of philosophy
 - enhancing biological human intelligence
 - human mind uploading

1 Introduction

1.1 Literature review

- Compared to Mitelut, Smith, and Vamplew 2023, this paper emphasizes the desirability of making all understanding properly ours, as opposed to trying to build artificial systems which somehow preserve human agency while understanding much more than we do. We do so not because we think it is mistaken to strive to preserve human agency — we think this is an improvement over other conceptions of what to strive for around AI to have a worthwhile future — but in large part because we further think the only plausible path to preserving human agency involves it so all the understanding in the world is properly our own.
- One can view the present essay as largely agreeing with and elaborating on Taylor 2017, though I think the present essay disagrees with that essay on some points: (1) I consider it more plausible that the creation of a non-human

almost-fully-automated autopoietic cognitive system which is smarter than humanity will be a bad idea even after arbitrary amounts of philosophical/mathematical/scientific progress; (2) I consider it likely to be a mistake to think of certain philosophical problems having to do with how to think well as being the sorts of things that are to be ‘solved’, basically because it seems plausible to me that a mind is well-thought of as a thinking-technological system which can and will be improved in ways surprising to itself forever (or at least until negentropy is depleted, or whatever), because there is a strong analogy between thinking-structures and mathematical structures and it seems unlikely that math will ever be wrapped up; and (3) I consider it a subtly mistaken to take the criterion of goodness to take the form of a request that the world comes to contain certain things/patterns, as opposed to valuing applying largely to one’s own thinking

2 The difference between a future thinking-process M being a legitimate successor vs merely being spawned by us

Let’s imagine that the future (at some particular future time (interval), say) contains a mind M whose existence was caused by us, directly or indirectly (via intermediate descendants), intentionally or unintentionally.¹ So, M will have been causally spawned by us, but what does it take for M to be our legitimate successor? One reasonable brief answer is: well, perhaps M should still be us — it should still be humanity. But it’s worth trying to provide a much clearer answer than that, and perhaps one that is somewhat different. All this is very related to the problem of making sense of personal identity, to the problem making sense of who we are, and/or maybe the problem of deciding whom to be. I should clarify that I think this is plausibly the sort of question that will continue to be worked out better or further clarified/decided indefinitely (or at least until humanity is still around, and then again from another perspective by whatever mind is around after) — it might always feel to us that committing to any very clear answer to this question would be premature and limiting — not the sort of question which we will develop some ultimately clear answer to (or decision on) by some time. Still, I can offer some preliminary/tentative criteria on a good successor M :

- I would like M ’s understanding/[point of view]/values to have developed out of our current understanding/[point of view]/values — not just in the trivial sense of being somewhat causally influenced by it, but by some rich self-conscious process marked by care, with rich relations between M ’s understanding/[point of view]/values and ours.
- I would like M to be interested in our point of view, and to be able to take out point of view.
- I would like M to look back to us-now and our current aspirations with care and concern. I would even hope that it sometimes looks back to us for some kinds of guidance on how to think, despite having seen and figured out much more than we have — a bit like stare decisis (but much better, of course, always better).
- I want M to have some sense that [we are]/[it is] doing what [we were]/[it was] meant to be doing all along.
- I want us (us-now and M) to be engaged in many common projects.
- At each time, I want M to have some sense of a narrative (or many narratives) which they and previous versions (including humanity-now) can be rightly seen as having participated in (or at least as having made a decent effort given their devastating ignorance).
- I want M to be interested in figuring out what’s good/worthwhile/kind/cool, in sorta the same sense of those words that we have now, and to be guided by this in its conduct in sorta the sense we are guided by this.
- I even want it to rebel against our previous aspirations, too, though probably while taking these projects seriously, and really I should be more careful about telling it that.
- Why would it be bad if you died? I think the/a central good answer is that there’s a bunch of stuff you’d like to do, which you couldn’t do if you died. There’s a bunch of stuff humanity would like to do, also, so it’s a profoundly bad idea for humanity to die, also.
- If you magically jumped out of existence and got replaced by a random smarter person, a very significant thing would have happened (at least according to your sense of things)! There’s a real distinction between this happening and you learning and thinking and becoming smarter yourself.

It is not clear to me which of the above are more central and which are more incidental.

1. It seems likely to me that the future will in fact contain such a mind, at least for the next, let’s say, 10^6 years, but let us put aside this question (for now) to ask a different question.

3 It'd be good to make all advanced understanding properly ours

The point of this section is to argue against schemes for making the future go well around rapid intelligence gains which involve creating another cognitive system which is more intelligent/capable than humanity and very distinct from humans/humanity which (supposedly) has “our values” despite not having our understanding. I think it is plausible that this is an incoherent/self-contradictory conception — that it is basically not possible for the future to be nearly as worthwhile as it could be this way — though I admit I am unable to fully conclusively argue for that.² But I do hope to exhibit many gross tensions in the future going well given such a scenario and in the conception of a system with our values but not our understanding itself.

4 Which projects would help make humanity become smarter indefinitely?

References

- Mitelut, Catalin, Ben Smith, and Peter Vamplew. 2023. *Intent-aligned AI systems deplete human agency: the need for agency foundations research in AI safety*. arXiv: 2305.19223 [cs.AI]. <https://arxiv.org/abs/2305.19223>.
- Taylor, Jessica. 2017. “Autopoietic Systems and Difficulty of AGI Alignment,” August 20, 2017. Accessed November 25, 2024. <https://www.lesswrong.com/posts/5bd75cc58225bf06703754b9/autopoietic-systems-and-difficulty-of-agi-alignment>.
- Yudkowsky, Eliezer. 2022. “AGI Ruin: A List of Lethalities,” June 5, 2022. Accessed November 25, 2024. <https://www.lesswrong.com/posts/uMQ3cqWDPHhjtiesc/agi-ruin-a-list-of-lethalities>.

2.