

Some suggestions for reforming the decision-making setup

July 2024

1 Instead of probabilities of harm, let's use expected utilities

I would suggest switching from talking about probabilities of harm to talking about expected utilities. Here's why:

- Talking about expected utilities is strictly more general. If one picks the utility to simply be 0 if harm happens and 1 otherwise, then for any event E , we have
- We can likely come up with much better choices for the utility function than the indicator function which corresponds to tracking the probability of harm.

2 Inaction should probably be just another action

3 We could just ask the system to make the future good, instead of asking it do some particular thing and also keeping the future good